

C Primer

Karl Hallsby

Last Edited: October 7, 2020

Contents

1	Introduction	1
1.1	Properties	1
1.1.1	Imperative	1
1.1.2	Procedural	1
1.1.3	Lexically Scoped	1
1.1.4	Statically Typed	1
1.1.5	Weakly Type Checked	1
2	Syntax	1
2.1	Primitive Types	2
2.1.1	Integer Type Prefixes	2
2.2	Basic Operators	2
2.2.1	Arithmetic Operators	2
2.2.2	Logical	3
2.2.3	Relational	3
2.2.4	Assignment	3
2.2.5	Conditional Operator	3
2.2.6	Bitwise Operators	3
2.3	Boolean Expressions	4
2.4	Control Flow	4
2.4.1	Branching	4
2.4.1.1	<code>if</code>	4
2.4.1.2	<code>if else</code>	5
2.4.1.3	<code>if else if else</code>	5
2.4.1.4	<code>switch case</code>	5
2.4.2	Repetition	5
2.4.2.1	<code>while</code>	6
2.4.2.2	<code>for</code>	6
2.4.2.3	<code>do while</code>	6
2.5	Variables	7
2.5.1	Visibility	7
2.5.1.1	Global Variables	7
2.5.1.2	<code>extern</code> Variables	7
2.5.2	Lifetime	8
2.6	Functions	8
2.6.0.1	Declaration	8
2.6.0.2	Definition	8
2.6.1	Passing Parameters	9
2.6.1.1	Pass-by-Value	9
2.6.1.2	Pass-by-Reference	9

3	Pointers	9
3.1	Pointer Syntax	9
3.1.1	Declaration	9
3.1.2	Getting an Address	9
3.1.3	Dereferencing	9
3.2	<i>Why</i> have Pointers?	9
3.3	Uninitialized Pointers	13
3.4	<code>NULL</code> Pointers	13
3.5	<code>void</code> Pointer	13
4	Arrays	13
5	Strings	15
5.1	String Utilities	15
6	Memory Management	17
6.1	<code>malloc</code>	17
6.2	<code>calloc</code>	17
6.3	<code>realloc</code>	17
6.4	<code>free</code>	17
7	Composite Data Types	18
7.1	<code>struct</code>	18
7.2	<code>union</code>	18
7.3	<code>enum</code>	18
7.4	<code>typedef</code>	18
8	Compilation	20
8.1	Stages	20
8.2	Makefiles	20
A	man Pages	20

1 Introduction

This document is intended for programmers that are newer to the C language and its facilities. This is meant as a quick, supplementary, reference document for these programmers. Many of the code examples in this document are either from IIT's CS 351 course, or Kernighan and Ritchie's Kernighan and Ritchie 1978 manual, *The C Programming Language*.

1.1 Properties

C is referred to as “low-level” today. That means there are relatively few abstractions and few “syntactic sugars” for expressing computations. This means when you write C, you can typically guess what the assembly would look like, which also lends itself to C's execution speed. However, this also means that you have **VERY FEW** built-in language protections for your computations. So, you open yourself up to a whole new class of problems when developing and writing your programs. The language will not protect you, but C compilers will typically attempt to throw warnings or errors about the most egregious errors you may write.

Here are some properties of C that you should know about.

1.1.1 Imperative

C is an imperative language, meaning you express computation as a series of steps. This is based off a finite-state based understanding of computation.

This stands in stark contrast to functional languages, which express computation as function applications to expressions.

1.1.2 Procedural

A procedural language allows you to organize repeated computations into logical blocks, typically referred to as functions or procedures.

1.1.3 Lexically Scoped

C is lexically scoped because variables inside of procedures cannot exist outside of their logical blocks.

This stands in contrast to Dynamically Scoped languages, where variables are sometimes available outside of the written scope for that variable. Bash is an example of a dynamically scoped language.

1.1.4 Statically Typed

C is statically typed, because the types of **ALL** expressions **MUST** be specified **BEFORE compilation**. This also means that when something is declared to be a certain type, it stays that way, unlike Python. This means that type checking happens during compilation, ensuring that all expressions have well-formed types.

To ensure program flexibility, we also introduce type-casting, where we either widen the type or narrow it. For example, taking an `int` and turning it into a `long` is a widening type cast, which are usually safe. This means that sometimes we must deal with type polymorphism, although C's handling of this class of problems is minimal and quite basic.

1.1.5 Weakly Type Checked

C is technically strongly typed on its operations, however, it becomes weakly typed because compilers cannot always ensure well-formedness of expressions when Pointers are used. Because C is weakly typed, you are not always guaranteed that when you access data that you are interpreting the bytes the right way.

C becomes weakly-typed because you can typecast pointers, pull values out of unions with different types, and in general do weird things with anything in memory. This also means that the compiler might not throw type-checking errors during the compilation phase of program development. Some of these pointer issues will only arise after running the program and ensuring that it is tested properly.

Some of these issues are illustrated in Listing 1.

2 Syntax

The C language helped define a whole class of syntax that is used by many programming languages today. C's syntactic decisions can be seen in Java, Rust, C++, C#, and many others.

```

1  #include <stdio.h>
2
3  int main(void) {
4  /* Types are implicitly converted */
5      char c = 0x41424344;
6      int i = 1.5;
7      unsigned int u = -1;
8      float f = 10;
9      double d = 2.5F; // Note: 'F' suffix for floating point literals
10
11     printf("c = '%c', i = %d, u = %u, f = %f, d = %f\n", c, i, u, f, d);
12
13     /* Typecasts can be used to force conversions */
14     int r1 = f/d;
15     int r2 = f / (int) d;
16
17     printf("r1 = %d, r2 = %d\n", r1, r2);
18     return 0;
19 }

```

```

1  $ ./a.out
2  c = 'D', i = 1, u = 4294967295, f =10.00000
3  r1 = 4, r2 = 5

```

Listing 1: Illustration of C's Weak Type Checking

2.1 Primitive Types

C has just 4 primitive types:

char: One byte integers (0–255), meant to represent ASCII characters.

int: Integers, which is defined to be *at least* 16 bits.

- Additional prefixes can be used to increase or decrease the range of the integer.
- These are shown in Section 2.1.1.

float: Single precision IEEE floating point number.

double: Double precision IEEE floating point number.

2.1.1 Integer Type Prefixes

signed: The default for integers, meaning you **do not** have to specify this. Can represent both negative and positive integers. The range for this is $-2^{\# \text{ bits}-1}$ to $2^{\# \text{ bits}-1}-1$

unsigned: , Can only represent 0 and positive integers. Its range is 0 to $2^{\# \text{ bits}}-1$

short: Tells the compiler that the integer must be at least 16 bits.

long: Tells the compiler that the integer must be at least 32 bits.

long long: Tells the compiler that the integer must be at least 64 bits.

2.2 Basic Operators

Operators perform some operation on expressions. This could be an arithmetic, a relational, logical, etc. operation.

2.2.1 Arithmetic Operators

These operators are for performing mathematical operations. These are well-defined for integers and floating-point numbers.

- + The addition operator. Works similarly for integers and floating-point numbers.

- The subtraction operator. Works similarly for integers and floating-point numbers.
- * The multiplication operator. Works similarly for integers and floating-point numbers.
- / The division operator. This returns the quotient of a division. This has a different result for integers and floating-point numbers.
 - Integers return the quotient of the division, but no fractional part.
 - Floating-points return the entire remainder of the division.
- % The modulo operator. Returns the remainder of a division operation when dividing integers. Note that this is only defined for integers.

2.2.2 Logical

Logical operators work with boolean values.

- ! The logical NOT operator.
- && The logical AND operator. Returns 1 if and only if the left and right expressions are **BOTH** 1 at the same time. Otherwise, 0 is returned.
- || The logical OR operator. Returns 0 if and only if both the left and right expressions are 0 at the same time. Otherwise, 1 is returned.

2.2.3 Relational

These relational operators are used to define a value in relation to another. Typically, these are used for boolean comparisons.

- == The equality operator. Returns 1 if and only if the two expressions are equal, 0 otherwise.
- != The inequality operator. Returns 0 if and only if the two expressions are not equal, 1 otherwise.
- > The greater-than operator. Returns 1 if and only if the left expression has a greater value than the right one.
- >= The greater-than-or-equal-to operator. Returns 1 if and only if the left expression has a greater value or equal value than the right one.
- < The less-than operator. Returns 1 if and only if the left expression has a lesser value than the right one.
- <= The less-than-or-equal-to operator. Returns 1 if and only if the left expression has a lesser value or equal value than the right one.

2.2.4 Assignment

Unlike many other languages, in C, the assignment operator = is also an expression. This means that when an assignment is performed, it also returns a value, in this case, it returns the value that was assigned to that particular name.

- = The assignment operator. Assigns a value to a given name. Returns the value of the assignment.
- += The add-and-assign operator. Takes the name on the left, adds the value on the right to the value on the left, and stores the result in the value on the left.
- *= The multiply-and-assign operator. Takes the name on the left, multiplies the value on the left by to the value on the right, and stores the result in the value on the left.

Only += and *= are shown, but there are similar ones defined for other Arithmetic Operators too.

2.2.5 Conditional Operator

There is only one conditional operator, sometimes called the ternary operation or conditional expression. It is defined like so: `bool ? true_exp : false_exp`.

2.2.6 Bitwise Operators

Bitwise operators work on the component bits of a number. This means they behave slightly differently than any other operator, but are not typically used in day-to-day calculations. Usually, they are used to efficiently work with memory.

- & Bitwise AND
- | Bitwise OR
- ^ Bitwise exclusive OR (XOR)

~ Bitwise negation, one's complement.

>> Bitwise SHIFT right

<< Bitwise SHIFT left

2.3 Boolean Expressions

Because C is so low-level, the concept of `true` and `false` are defined as integers.

0 `false`.

1 `true`. Technically, any non-zero value is considered `true`.

```
1  #include <stdio.h>
2
3  int main(void) {
4      printf("%d\n", !(0));           // 1
5      printf("%d\n", 0 || 2);         // 1
6      printf("%d\n", 3 && 0 && 6);     // 0
7      printf("%d\n", !(1234));        // 0
8      printf("%d\n", !(-1020));       // 1
9
10     return 0;
11 }
```

```
1  $ ./a.out
2  1
3  1
4  0
5  0
6  1
```

Listing 2: Logical Operators

2.4 Control Flow

Control flow is the idea of changing the direction a program executes based on some predicate. Whether this change in direction is to a new direction is because of a change in state, or because something must be repeated is irrelevant.

2.4.1 Branching

Branching has to deal with the changing of a program's control flow based on a predicate to perform some separate actions based on the state of the predicate. There are 4 types for this:

if The most basic change in flow control. If the predicate provided is `true`-thy, performs an action, then returns to normal.

if else If the predicate is `true`-thy, then perform some action, if the predicate is `false`, then perform some other action.

if else if else If the predicate in the first `if` is `true`-thy, then that branch is taken. If the first predicate is `false`, then the `else if`'s predicate is checked for truth. The the `else if`'s predicate is `false`, then execution continues through any other `else if`s that may be present. If none of the `if` or `else if`s' predicates were `true`-thy, then the `else` is taken.

switch case The `switch-case` statement allows you to choose from many different paths based on the **VALUE** of some expression.

2.4.1.1 if The `if` statement has a very basic syntax, shown below in Listing 3. Typically, this is only used when there needs to be a small change in the state of the program based off the predicate's value.

```

1  if (logical-predicate) {
2      predicate-true-clause;
3  }
4
5  /* Note that the ( and ) MUST be present in the statement to denote the predicate.
6   * { and } can be omitted ONLY IF the clause has a SINGLE statement in it. */

```

Listing 3: `if` Statement Syntax

2.4.1.2 if else The `if else` statement is used to perform two distinct actions based on the predicate's value. The syntax of this is shown in Listing 4.

```

1  if (logical-predicate) {
2      predicate-true-clause;
3  } else {
4      predicate-false-clause;
5  }
6
7  /* Note that the ( and ) MUST be present in the statement to denote the predicate.
8   * { and } can be omitted ONLY IF the clause has a SINGLE statement in it. */

```

Listing 4: `if else` Statement Syntax

2.4.1.3 if else if else The `if else if else` is used to choose between n options. However, this comes with the downside that **EACH** potential path's predicate **MUST** be evaluated before any action can occur. The syntax for this is shown in Listing 5.

```

1  if (logical-predicate) {
2      this-predicate-true-clause;
3  } else if (logical-predicate) {
4      this-predicate-true-clause;
5  } else {
6      all-predicates-false-clause;
7  }
8
9  /* Note that the ( and ) MUST be present in the statement to denote the predicate.
10 * { and } can be omitted ONLY IF the clause has a SINGLE statement in it. */

```

Listing 5: `if else if else` Statement Syntax

2.4.1.4 switch case The `switch case` statement is used to choose between n different options. However, unlike the `if else if else`, the expression is only evaluated the once, and the branch is then jumped to in constant time, making this a better option for making one decision of many. The syntax for this is shown in Listing 6.

2.4.2 Repetition

Here, we want to perform some action a number of times. There are 3 structures for performing a repeating action:

while Loop This is generally used when the number of repetitions is unknown or uncountable at any given point in time. It continues to repeat until the predicate ceases to be `true`-thy or an explicit `break` occurs.

```

1  switch (predicate) {
2  case 1:
3      break;
4  case 2:
5      expression1;
6  case 3: {
7      statement1;
8      statement2;
9  }
10 default:
11     break;
12 }
13
14 /* The ( and ) are MANDATORY for the switch to occur. In addition, the outer { and }
15  * are MANDATORY for each of the cases.
16  * If there will be many STATEMENTS in a case, then the body of the case MUST be
17  * surrounded by { and }.
18  * In addition, each case will "fall-through" by default, meaning the body of an
19  * case will also execute the bodies of lower cases. The only way to prevent this
20  * action is to include a break in the body of the case. */

```

Listing 6: `switch case` Statement Syntax

for Loop The `for` loop is used when the number of repetitions is both known and countable.

do while Loop The `do while` loop is the same as the `while` loop, but the difference is that the body of the loop is executed **ONCE before** evaluating the predicate.

2.4.2.1 while The `while` loop is typically used when the number of repetitions is unknown or uncountable. This allows for easy definition of infinite loops, allowing for a program to continue execution until some condition has been met. The syntax for this is shown in Listing 7.

```

1  while (condition) {
2      statements;
3  }
4
5  /* The body of the while-loop will continue to execute UNTIL condition becomes
6  * false.
7  * The ( and ) are NEEDED for the predicate. In addition, the { and } are
8  * NEEDED if there is more than one statement in the body. */

```

Listing 7: `while` Loop Syntax

2.4.2.2 for The `for` loop is used when the number of repetitions is both countable. This type of loop can be modelled by the `while` as well, but this helps ensure that all the components of the loop are present, helping prevent infinite loops. The syntax of this is shown in Listing 8.

2.4.2.3 do while The `do while` loop performs some action an uncountable or unknown number of times, **AT LEAST** once. This is typically used when the action to be repeated can be repeated any number of times, but must happen at least once. The syntax for this is shown in Listing 9.

```

1  for(init value; predicate; per-loop change) {
2      statement;
3  }
4
5  /* The body of the ( and ) part must be written that way, and MUST be in the
6  * parentheses.
7  * The { and } are MANDATORY if there is MORE than 1 statement in the for loop.
8  * The initial value MUST be countably changeable, meaning an integer or
9  * something else.*/

```

Listing 8: **for** Loop Syntax

```

1  do {
2      statement;
3  }
4  while(condition);
5
6  /* The { and } SHOULD ALWAYS be included for the body of a do-while loop, although
7  * the language and compiler do NOT require them. However, the ( and ) ARE required
8  * to surround the condition.
9  * In addition, the ending semicolon ; at the end of the while IS mandatory. */

```

Listing 9: **do while** Loop Syntax

2.5 Variables

Because C is a statically typed language, the type of every expression **MUST** be known before or during compilation. In addition, C compilers do not have any type inferencing, so you **MUST** explicitly tell the compiler the type of your variables. This means that you **MUST** declare before use. It is important to note that the declaration implicitly allocates storage for the data that will be stored.

One thing that will come up throughout this section is the concept of aliveness and scope. It is important to note that variables do **not** have to be in-scope to be alive, and vice versa.

2.5.1 Visibility

Visibility or *scope* is where a symbol can be seen from. If the symbol cannot be seen, then it cannot be used in any way

In addition, we need to ask *how* we can refer to the symbol. This includes what kind of identifiers/modifiers/namespacing is needed to identify the symbol in question.

2.5.1.1 Global Variables They **MUST** be declared outside any function. These are not deallocated **AT ANY TIME** during a program's execution, as they are always in-scope, however the variable may not be alive. Additionally, these are **ALWAYS** available, until the program terminates.

Using global variables is typically bad practice as this can introduce weird and hard-to-debug errors into a program. So, most variables that you will use will be local variables. Local variables are limited to the scope they were created within. Typically, the scope is a function, but can be an **if**, a **while**, etc.

2.5.1.2 extern Variables **extern** is used to denote a variable that is external to this program. This means the variable in question is a global variable in another file

The opposite of the **extern** keyword is the **static** keyword, which limits the scope of a symbol to the file it is declared in. In addition, when the variable is declared to be **static**, the value lasts throughout the program's execution, ensuring the variable is always available.

2.5.2 Lifetime

The lifetime of a variable asks how long a variable has storage allocated to it. In C, this is distinctly different than a symbol not being visible. One example of this is: a pointer can have the memory underneath it deallocated, ending the lifetime of the pointer, but keeping the pointer in-scope.

2.6 Functions

Functions are the highest form of modularity present in the language. Here, the procedural aspect of the language comes into play, and becomes useful. Repeated computations can be expressed as a function, allowing for easy code reuse and modularization.

Because C is so low-level, it has the unique distinction of requiring the programmer to separate the declaration of a function from its definition.

2.6.0.1 Declaration A declaration of a function simply announces to the compiler that a function with those input parameters and output values will be defined. However, a declaration says nothing about **HOW** the function will be defined, only that one such function will exist.

In C programming, function declarations (and sometimes some variable definitions) are placed in *header* files, that end with the *.h* extension. Sometimes, the function declarations here are called *function prototypes*.

An example of this is shown in Listing 10.

```
1 unsigned long hash(char *str);
2 hashtable_t *make_hashtable(unsigned long size);
3 void ht_put(hashtable_t *ht, char *key, void *val);
```

Listing 10: Declaration of `hashtable.h`

2.6.0.2 Definition The definition of a function is the actual implementation of a function. These are typically done in *.c* files.

An example of this is shown in Listing 11.

```
1 #include "hashtable.h"
2
3 unsigned long hash(char *str) {
4     unsigned long hash = 5381;
5     int c;
6     while ((c = *str++))
7         hash = ((hash << 5) + hash) + c;
8     return 0;
9 }
```

```
1 #include "hashtable.h"
2
3 int main(int argc, char *argv[]) {
4     hashtable_t *ht;
5     ht = make_hashtable(atoi(argv[1]));
6     return 0;
7 }
```

Listing 11: Definition of `hashtable.c`

2.6.1 Passing Parameters

The act of passing parameters is key to actually using functions in our code. There are two ways to pass parameters through to functions in C.

1. Pass-by-Value
2. Pass-by-Reference

2.6.1.1 Pass-by-Value Passing by value involves making a copy of the value passed to a function and giving it a new name. Since this is a **COPY**, it means **ANY** modifications to the copy **DO NOT** affect the original.

This is useful for small pieces of information, but the problem is when the size of the information increases. This is because the entire contents of a thing must be copied before running the function. This is where the other form of parameter passing comes into play.

2.6.1.2 Pass-by-Reference Passing by reference involves giving the called function a pointer (reference) to the data to manipulate. This is extremely efficient for very large data structures, because only the pointer to the data needs to be copied around. This is particularly valuable for large `struct` s, arrays, etc.

3 Pointers

This should technically go in Section 2, but pointers deserve their own section.

Defn 1 (Pointer). A *pointer* is a variable declared to store a memory address. With this memory address, we can refer to data in-memory. The size of the pointer is determined by the architecture of the CPU.

A pointer is designated by its **DECLARED** type, **NOT** its contents. This allows the data the pointer points to to be re-interpreted based on the declared type of the pointer. This is shown in Listing 12.

3.1 Pointer Syntax

. The syntax that Pointers use can sometimes be confusing for new programmers. So, we will break down each portion of a pointer and its usage to more fully understand them.

3.1.1 Declaration

Pointers are declared using the same type name as the data type they will store. In addition, you **MUST** add a `*` to the declaration. The placement of this symbol does not matter, but for clarity, most programmers put it on the variable name. However, it can be attached to the type as well.

Listing 13 shows how to do this.

3.1.2 Getting an Address

If you already have a name that points to the actual value, and you want the address of that value, you can use the `&` unary operator. Its usage is shown in Listing 14.

3.1.3 Dereferencing

In a major point of confusion for the C language, the `*` operator is **also** used to **dereference** the pointer! An example of declaring, getting the address of a variable, and the dereferencing of pointers is shown in Listing 15.

3.2 Why have Pointers?

Pointers are partly a cross-over from assembly, as C is really just a thin wrapper around that. However, it does also allow direct access to memory and allow *us* to make many optimizations we could not make otherwise. This means you must **manage memory yourself**.

- In C, everything is *ALWAYS* Pass-by-Value.
- This even happens on composite data structures, like `struct` s.
- Pointers enable us to perform *actions at a distance*, essentially allow us to get Pass-by-Reference without explicitly allowing that.
- This also allows functions very deep in the call stack to affect variables earlier in the stack.

One good reason is illustrated in Listing 16.

```

1  #include <stdio.h>
2  #include <stdlib.h>
3
4  int main(void) {
5      int *ip;
6      char *cp;
7      float *fp;
8
9      ip = malloc(1); /* Allocates one byte. */
10
11     *ip = 63; /* Store the integer 63 in the memory location that `ip' points to. */
12
13     /* Let cp and fp point to the same memory location as ip. */
14     cp = (char*) ip;
15     fp = (float*) ip;
16
17     /* Print the value that all 3 pointers point to in memory. */
18     printf("ip: %d\n", *ip);
19     printf("cp: %c\n", *cp);
20     printf("fp: %f\n", *fp);
21
22     return 0;
23 }

```

```

1  $ ./a.out
2  ip: 63
3  cp: ?
4  fp: 0.000000

```

Listing 12: Pointers Reinterpret Data

```

1  /* This struct has no bearing on how the declaration of a pointer works.
2   * I just wanted to give a thorough example of how pointer declaration works. */
3  struct student {
4      char *name;
5      float grade;
6      char letter_grade;
7  };
8
9  int main(int argc, char *argv[])
10 {
11     /* type *var_name; */
12     int *ip;
13     char *cp;
14     struct student *sp;
15     return 0;
16 }

```

Listing 13: Pointer Declaration

```

1  #include <stdio.h>
2
3  int main(void) {
4      int x;
5      int *xp;
6
7      x = 15; /* Assign 15 to the location of x. */
8
9      xp = &x; /* Get the address of x, and store in xp. */
10
11     printf("x: %d\n", x);
12     printf("*xp: %d\n", *xp);
13     printf("xp: %p\n", xp);
14
15     return 0;
16 }

```

```

1  $ ./a.out
2  x: 15
3  *xp: 15
4  xp: 0x7ffc4868a204

```

Listing 14: Address & Operator

```

1  #include <stdio.h>
2
3  int main(void) {
4      int i, j, *p, *q;
5
6      i = 10;      /* i is defined to have the value 10 */
7      p = &j;      /* j is uninitialized still, put j's address in p */
8      q = p;      /* Set the pointer q to the address in p, which is j's address */
9      *q = i;      /* Store i in the location q POINTS to, which is j right now */
10     *p = *q * 2; /* Multiply the value q POINTS to by 2 and store in the location POINTED to by p
11     ↪ */
12     printf("i=%d, j=%d, *p=%d, *q=%d\n", i, j, *p, *q);
13     return 0;
14 }

```

```

1  $ ./a.out
2  i=10, j=20, *p=20, *q=20

```

Listing 15: Dereferencing Pointers

```

1  void swap(int x, int y);
2  void p_swap(int *p, int *q);
3
4  int main() {
5      int a = 5, b = 10;
6      swap(a, b);
7      /* want a == 10, b == 5
8      /* swap() doesn't do that though. */
9      p_swap(&a, &b);
10     /* Pass memory addresses into p_swap, allowing for direct access.*/
11     return 0;
12 }
13
14 /* This uses call-by-value and does the swap in function-local variables.
15 /* This means that the original a and b would NOT be changed. */
16 void swap(int x, int y) {
17     int tmp = x;
18     x = y;
19     y = tmp;
20     return;
21 }
22
23 /* This uses call-by-reference and does the swap on the original memory locations.
24 /* This means that the original a and b WOULD be changed. */
25 void p_swap(int *p, int *q) {
26     int tmp = *p;
27     *p = *q;
28     *q = tmp;
29     return;
30 }

```

Listing 16: Swap with Pointers

3.3 Uninitialized Pointers

Even though a Pointer can only store memory addresses, it behaves exactly like a regular variable for most puposes.

- When a Pointer is declared, they are given their underlying storage, which **will contain garbage!**
- If you dereference this pointer, you dereference garbage, which is undefined behavior.
 - If you're lucky, this is will result in a **SEGFAULT**, terminating program execution.
 - If you're unlucky, unknown results may happen, and you might never know about it.
 - You are assured that your program will not exceed its own memory space though, which means your program cannot crash another.

Listing 17 gives an example of what could happen if you don't initialize your Pointers.

```
1  #include <stdio.h>
2
3  int main(void) {
4      int *ip;
5
6      /* Don't assign anything to ip, which means it holds garbage. */
7      printf("%d\n", *ip);
8
9      return 0;
10 }
```

```
1  $ ./a.out
2  [1] 29456 segmentation fault (core dumped) ./a.out
```

Listing 17: Uninitialized Pointers

3.4 NULL Pointers

The **NULL** Pointer is **NEVER** returned by the `&` operator. It is usually a smart idea to define pointers when you declare them, and if you don't immediately define them with a usable value, to define them with **NULL**. This helps prevent many runtime issues that the compiler cannot check for. Listing 18 gives an example of using the **NULL** pointer this way.

- The **NULL** pointers is safe to use as a predicate, in **if** statements, for example. This usage is shown in Listing 19.
- Written as 0 or **NULL** in *pointer context*
 - Typically `#define NULL 0`

3.5 void Pointer

This is a pointer type that says there is no type. It is compatible with **EVERY** other Pointer type, but it **CANNOT** be used directly. The **void** pointer **MUST** be typecast before being used in **ANY** way.

4 Arrays

An Array is the only way to store multiple items under a single name.

Defn 2 (Array). A *array* is a contiguous block of memory. Arrays are indexed items, meaning they use a number to identify each item.

```
1  #include <stdio.h>
2
3  int main(void) {
4      int *ip = NULL;
5
6      /* Don't assign anything to ip, which means it holds garbage. */
7      printf("%d\n", *ip);
8
9      return 0;
10 }
```

```
1  $ ./a.out
2  [1] 29456 segmentation fault (core dumped) ./a.out
```

Listing 18: `NULL` Pointer as Checker

```
1  #define NULL 0
2
3  int main() {
4      int i = 0;
5      int *p = NULL;
6
7      if (p) {
8          /* Looks safe to dereference */
9      }
10     return 0;
11 }
```

Listing 19: `NULL` Pointer as Sentinel

Arrays are declared with the following syntax `type arr_name[size]` . This declaration also implicitly allocates space for its storage. This allocates the space from the function **stack** *NOT* the heap.

There is **NO** metadata about the array, such as its length. This means:

- **NO** implicit size
- **NO** bounds checking

All of the ways to statically declare, allocate, and define an array are shown in Listing 20.

```
1  int main(void) {
2      int i_arr[10]; /* Array of 10 ints, 40 bytes */
3      char c_arr[80]; /* Array of 80 characters, 80 bytes */
4      char td_arr[24][80]; /* 2-D Array, 24 x 80 x 1 bytes*/
5      int *ip_arr[10]; /* Array of 10 pointers which point to ints, 40 or 80 bytes */
6
7      /* Dimension inferred if initialized when declaring. */
8      short grades[] = { 75, 90, 85, 100 };
9
10     /* Can omit the first dimension, as partial initialization is allowed. */
11     int sparse[][10] = { { 5, 3, 2},
12                          { 8, 10 },
13                          { 2 } };
14
15     /* If partially initialized, remaining components are 0 */
16     int zeros[1000] = { 0 }; /* Initialize all values to 0 */
17
18     /* Can also use designated initializers for specific indices */
19     int nifty[100] = { [0] = 0, /* Element zero has value 0*/
20                      [99] = 1000, /* Element 99 has value 1000 */
21                      [49] = 250 };
22
23     return 0;
24 }
```

Listing 20: Arrays, their Declaration and Definition

The syntax of an Array actually is syntactic sugar for Pointer arithmetic. The name we give to the array is actually the name of the **Pointer** that points to the beginning of the array, the zeroth element in the array. This is shown in Listing 21.

Because arrays are fancy pointers, you can also typecast arrays.

It is important to note that arrays that are present on the stack of a program **MUST** be of fixed size. This is because the compiler needs to know how much space to allocate for each procedure in the program. However, if we want a dynamically sized array, we will need to turn to the heap. Using the heap is discussed in Section 6. To access the heap from our program, we need a Pointer to point to the array on the heap. By doing this, we still have a fixed-size stack, because pointers are of fixed size, but we allow use of dynamically sized data structures.

5 Strings

In C, strings are actually character arrays. These arrays are terminated with the null character, `\0` ; it's numerical value is 0. This is shown in Listing 22.

`printf` treats strings as a character array terminated by a null character.

5.1 String Utilities

All the functions shown below string are from `<string.h>`.

`char *strcpy(char *dest, const char *src)` Copy characters from source to destination array, including the `\0` .

```
1  #include <stdio.h>
2
3  int main(void) {
4      int x[] = { 1, 2, 3, 4, 5 };
5
6      printf("3rd element in x, using array syntax: %d\n", x[2]);
7      /* Because x is an int, the +2 actually performs the + (2 * sizeof(int)) operation. */
8      printf("3rd element in x, using pointer syntax: %d\n", *(x + 2));
9
10     return 0;
11 }
```

```
1  $ ./a.out
2  3rd element in x, using array syntax: 3
3  3rd element in x, using pointer syntax: 3
```

Listing 21: Array-Pointer Similarity

```
1  #include <stdio.h>
2
3  int main() {
4      char *str = "hello world!"; /* 12 characters + 1 null character. */
5      str[12] = 10; /* Set the null terminator to 10. */
6      printf("%s", str);
7      return 0;
8  }
```

```
1  $ ./a.out
2  hello world!
```

Listing 22: String/Character Array

`char *strcat(char *dest, const char *src)` Concatenates 2 strings, removing null bytes and adding one to the very end of the destination string.

`int strcmp(const char *s1, const char *s2)` Compares strings byte-by-byte, returning whichever string is greater than the other, as determined by the component characters' ASCII code.

`size_t strlen(const char *s)` Finds length of string by incrementing a counter until it finds the null character.

`void *memcpy(void *dest, const void *src, size_t n)` Copies the contents of one memory location to another.

`void* memmove(void *dest, const void *src, size_t n)` Moves the contents of one memory location to another.

6 Memory Management

Because C is a language that does not provide many abstractions, it also requires the programmer to remember and manage their memory usage. So, **YOU** must be the one to manage the memory, there is **NO** built-in garbage collector for you to use.

Memory allocation is done on the heap of the program's execution space in memory. When you allocate memory in your program, you are actually requesting the operating system to give you the memory you want.

6.1 malloc

This is the simplest function of all possible memory allocation functions. `malloc`:

- Takes one argument:
 1. The number of bytes to allocate.
 - Returns a **POINTER** to the front of the allocated memory.
- `malloc` ***DOES NOT*** initialize memory, so it will be garbage.

6.2 calloc

This is quite similar to `malloc`. `calloc`:

- Takes 2 arguments:
 1. The number of spaces to allocate, for example the number of elements in an array.
 2. The number of bytes to allocate, for the type being stored.
 - Returns a **POINTER** to the front of the allocated memory.
- `calloc` ***ZEROS*** memory, so this does have a slight performance penalty.

6.3 realloc

`realloc` is used to **REALLOCATE** an existing memory location.

- Takes 2 arguments:
 1. The pointer to the memory location previously allocated with either `malloc` or `calloc`.
 2. The amount of memory to reallocate, in bytes.
- If the `NULL` pointer is passed to `realloc`, it will behave exactly like `malloc`.
- Returns a **POINTER** to the front of the reallocated memory

6.4 free

`free` is used to free memory that was previously allocated, removing from the programming space entirely.

- Takes 1 argument:
 1. A pointer to the memory to be deallocated.
- Returns `void`.

7 Composite Data Types

This is *similar* to objects in Object-Oriented Programming, or the `datatype` in Haskell.

7.1 `struct`

The `struct` keyword allows us to put multiple separate data types together and refer to them by their field name. You access the fields by using the `.` operator. An example of this is shown in Listing 23.

```
1  /* Type Definition */
2  struct point {
3      int x;
4      int y;
5  }; /* Ending ';' is REQUIRED. */
6
7  /* point declaration and allocation. */
8  struct point pt;
9
10 /* Pointer to a point struct */
11 struct point *pp1;
12
13 int main(void) {
14     pt.x = 10;
15     pt.y = -5;
16
17     struct point pt2 = {.x = 8, .y = 13 }; /* Declaration and Initialization. */
18
19     struct point *pp;
20     pp = &pt;
21
22     /* To access a field in pp, the 2 syntaxes are equivalent. */
23     (*pp).x;
24     pp->x;
25
26     return 0;
27 }
```

Listing 23: `struct` Usage

7.2 `union`

The `union` keyword allows us to define a single type that can be of one type from many. However, these are **NOT** like Haskell's `datatype` or Rust's union, in that we do not have type protections about accessing this data. Listing 24 gives an example of this.

7.3 `enum`

The `enum` allows us to have a set of mutually exclusive options among many. These are technically backed by an array of integers. However, this massively helps legibility, readability, and writability than its backing structure. An example of this datatype is shown in Listing 25.

7.4 `typedef`

The `typedef` keyword allows us to change the name of types. An example of this is shown in Listing 26.

```
1  #include <stdio.h>
2
3  union random {
4      int i;
5      float f;
6      char c;
7  }; /* Ending ';' is REQUIRED. */
8
9  int main(void) {
10     union random test;
11
12     test.i = 63;
13
14     /* This is allowed, and the program will print the corresponding character. */
15     printf("%c", test.c);
16
17     return 0;
18 }
```

Listing 24: `union` Usage

```
1  enum stoplight_t {
2      RED,
3      YELLOW,
4      GREEN
5  }; /* Ending ; IS mandatory. */
6
7  int main(void) {
8      enum stoplight_t state_31;
9
10     /* This is valid code. */
11     state_31 = GREEN;
12
13     return 0;
14 }
```

Listing 25: `enum` Usage

```
1  /* typedef oldname newname; */
2  typedef int int_t; /* int and int_t are aliases to the same type. */
```

Listing 26: `typedef` Usage

8 Compilation

C, along with C++ can be quite painful to compile for larger projects. You could manually compile every `.c` file with `gcc`. Makefiles help us manage this.

8.1 Stages

1. Preprocessing
 - Preprocessor directives starting with `#`
 - Text substitution
 - Macros
 - Conditional compilation
 - Performs complete textual substitution behind the scenes
2. Compile
 - From source language to object code/binary
3. Link
 - Put inter-related object codes together
 - Resolve calls/references and definitions
 - Put absolute/relative addresses into the binary for the `=call=` instruction
 - Want to support /selective/ public APIs
 - Don't always want to allow linking a call to a definition

8.2 Makefiles

Makefiles allow for:

- Incremental compilation
- Automated compilation

These are written as a list of targets, prerequisites, and directives. In addition, they can include variables to simplify typing things out. They also support arbitrarily long lists, file globbing (like regular expressions), and a couple of other things. An example **Makefile** is shown in Listing 27.

```
1 CC = gcc
2 CCFLAGS = -Wall -g
3
4 all: prerequisite1 prerequisite2
5     $(CC) prerequisite1 prerequisite2
6
7 target: prerequisite
8     directive
```

Listing 27: Example **Makefile**

A man Pages

The **man** pages, standing for manual pages, are reference documents for using these APIs. If you are curious about a function, you can just type `man func-name`.

These will give detailed information about the function, and possibly other, related, functions.

There are several sections of **man** pages, listed below:

1. General Commands
2. System Calls
3. Library functions, in particular the C standard library
4. Special files (usually devices, those found in `/dev`) and drivers
5. File formats and conventions
6. Games and screensavers
7. Miscellanea
8. System administration commands and daemons

References

- [KR78] Brian W. Kernighan and Dennis M. Ritchie. *The C Programming Language*. English. Second. Prentice Hall, Feb. 22, 1978. ISBN: 9780131101630.